

```
import pandas as pd  
import numpy as np
```

```
df = pd.read_csv("2023_Taxi_5M.csv")
```

```
/tmp/ipython-input-729068485.py:1: DtypeWarning: Columns (4,10,16) have mixed t  
df = pd.read_csv("2023_Taxi_5M.csv")
```

```
df.head()
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count	trip
0	2	01/01/2023 12:32:10 AM	01/01/2023 12:40:36 AM		1.0
1	2	01/01/2023 12:39:42 AM	01/01/2023 12:50:36 AM		1.0
2	2	01/01/2023 12:09:29 AM	01/01/2023 12:29:23 AM		2.0
3	2	01/01/2023 12:45:11 AM	01/01/2023 01:07:39 AM		1.0
4	1	01/01/2023 12:51:45 AM	01/01/2023 12:58:18 AM		1.0

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4278823 entries, 0 to 4278822  
Data columns (total 19 columns):  
 #   Column           Dtype  
---  --  
 0   VendorID        int64  
 1   tpep_pickup_datetime  object  
 2   tpep_dropoff_datetime  object  
 3   passenger_count    float64  
 4   trip_distance     object  
 5   RatecodeID        float64  
 6   store_and_fwd_flag object  
 7   PULocationID     int64  
 8   DOLocationID     int64  
 9   payment_type      float64  
 10  fare_amount       object  
 11  extra             float64  
 12  mta_tax            float64  
 13  tip_amount         float64  
 14  tolls_amount       float64  
 15  improvement_surcharge float64  
 16  total_amount       object  
 17  congestion_surcharge float64  
 18  airport_fee        float64  
dtypes: float64(10), int64(3), object(6)  
memory usage: 620.3+ MB
```

```
df.isna().sum()
```

	0
VendorID	0
tpep_pickup_datetime	0
tpep_dropoff_datetime	0
passenger_count	130107
trip_distance	0
RatecodeID	130107
store_and_fwd_flag	130107
PULocationID	0
DOLocationID	0
payment_type	1
fare_amount	1
extra	1
mta_tax	1
tip_amount	1
tolls_amount	1
improvement_surcharge	1
total_amount	1
congestion_surcharge	130108
airport_fee	130108

dtype: int64

```
df['airport_fee'] = df['airport_fee'].fillna(0)
df['congestion_surcharge'] = df['congestion_surcharge'].fillna(0)
```

```
print(df['passenger_count'].mean())
```

1.367321118148362

```
df['passenger_count'] = df['passenger_count'].fillna(1)
```

```
df = df.drop(columns=['store_and_fwd_flag'])
df['RatecodeID'] = df['RatecodeID'].fillna(99)
```

```
df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])
```

```
/tmp/ipython-input-1208249868.py:1: UserWarning: Could not infer format, so each column was converted to category type
  df['tpep_pickup_datetime'] = pd.to_datetime(df['tpep_pickup_datetime'])
/tmp/ipython-input-1208249868.py:2: UserWarning: Could not infer format, so each column was converted to category type
  df['tpep_dropoff_datetime'] = pd.to_datetime(df['tpep_dropoff_datetime'])
```

```
df.select_dtypes(include=['number']).agg(['min', 'max', 'mean']).T
```

	min	max	mean
VendorID	1.00	6.00	1.738111
passenger_count	0.00	9.00	1.356152
RatecodeID	1.00	99.00	4.580434
PULocationID	1.00	265.00	165.163665
DOLocationID	1.00	265.00	163.940594
payment_type	0.00	4.00	1.188986
extra	-7.50	96.38	1.568548
mta_tax	-0.50	5.75	0.486007
tip_amount	-81.00	480.10	3.518550
tolls_amount	-49.85	196.99	0.588766
improvement_surcharge	-1.00	1.00	0.980018
congestion_surcharge	-2.50	2.75	2.195470
airport_fee	-1.75	1.75	0.135036

```
df['trip_distance'] = df['trip_distance'].str.replace(',', '', regex=False)

df['trip_distance'] = df['trip_distance'].astype(float)
```

```
float_cols = [
    'passenger_count',
    'trip_distance',
    'RatecodeID',
    'fare_amount',
    'extra',
    'mta_tax',
    'tip_amount',
    'tolls_amount',
    'improvement_surcharge',
    'total_amount',
    'congestion_surcharge',
    'airport_fee'
]

for col in float_cols:
    if col in df.columns:
```

```

df[col] = (
    df[col]
    .astype(str)
    .str.replace(',', '', regex=False)
)
df[col] = pd.to_numeric(df[col], errors='coerce')

```

```
df.select_dtypes(include=['number']).agg(['min', 'max', 'mean']).T
```

	min	max	mean
VendorID	1.00	6.00	1.738111
passenger_count	0.00	9.00	1.356152
trip_distance	0.00	335004.33	5.948710
RatecodeID	1.00	99.00	4.580434
PULocationID	1.00	265.00	165.163665
DOLocationID	1.00	265.00	163.940594
payment_type	0.00	4.00	1.188986
fare_amount	-900.00	386983.63	19.598446
extra	-7.50	96.38	1.568548
mta_tax	-0.50	5.75	0.486007
tip_amount	-81.00	480.10	3.518550
tolls_amount	-49.85	196.99	0.588766
improvement_surcharge	-1.00	1.00	0.980018
total_amount	-901.00	386987.63	28.538945
congestion_surcharge	-2.50	2.75	2.195470
airport_fee	-1.75	1.75	0.135036

```

df = df[df['trip_distance'] > 0.5]
df = df[df['trip_distance'] < 100]

df = df[~df.select_dtypes('number').lt(0).any(axis=1)]

```

```
df.select_dtypes(include=['number']).agg(['min', 'max', 'mean']).T
```

	min	max	mean
VendorID	1.00	6.00	1.747906
passenger_count	0.00	9.00	1.337836
trip_distance	0.51	99.76	3.686328
RatecodeID	1.00	99.00	7.651442
PULocationID	1.00	265.00	164.975514
DOLocationID	1.00	265.00	163.595329
payment_type	0.00	4.00	1.115288
fare_amount	0.00	1025.00	20.360364
extra	0.00	14.25	1.610381
mta_tax	0.00	5.75	0.496961
tip_amount	0.00	280.00	3.681055
tolls_amount	0.00	94.75	0.620149
improvement_surcharge	0.00	1.00	0.999241
total_amount	0.00	1026.50	29.641226
congestion_surcharge	0.00	2.75	2.179459
airport_fee	0.00	1.75	0.136827

```
df = df[df['total_amount'] > 3.7]

df.select_dtypes(include=['number']).agg(['min', 'max', 'mean']).T
```

	min	max	mean
VendorID	1.00	6.00	1.747953
passenger_count	0.00	9.00	1.337851
trip_distance	0.51	99.76	3.686093
RatecodeID	1.00	99.00	7.649950
PULocationID	1.00	265.00	164.976648
DOLocationID	1.00	265.00	163.595775
payment_type	0.00	4.00	1.115181
fare_amount	0.00	1025.00	20.362166
extra	0.00	14.25	1.610515
mta_tax	0.00	5.75	0.496997
tip_amount	0.00	280.00	3.681384
tolls_amount	0.00	94.75	0.620204
improvement_surcharge	0.00	1.00	0.999271
total_amount	3.98	1026.50	29.643776
congestion_surcharge	0.00	2.75	2.179649
airport_fee	0.00	1.75	0.136829

```
df = df[df['passenger_count'] >= 1]
```

```
df.select_dtypes(include=['number']).agg(['min', 'max', 'mean', 'median', 'std'])
```

	min	max	mean	median	std
VendorID	1.00	6.00	1.759152	2.00	0.441703
passenger_count	1.00	9.00	1.357890	1.00	0.856223
trip_distance	0.51	99.76	3.695883	1.94	4.585362
RatecodeID	1.00	99.00	7.748972	1.00	24.706680
PULocationID	1.00	265.00	164.944502	162.00	64.111443
DOLocationID	1.00	265.00	163.578450	162.00	69.923377
payment_type	0.00	4.00	1.113693	1.00	0.514395
fare_amount	0.00	1025.00	20.398935	14.20	17.795989
extra	0.00	14.25	1.583503	1.00	1.840527
mta_tax	0.00	5.75	0.496983	0.50	0.039211
tip_amount	0.00	280.00	3.688199	3.00	4.006492
tolls_amount	0.00	87.50	0.623581	0.00	2.208256
improvement_surcharge	0.00	1.00	0.999264	1.00	0.022974
total_amount	3.98	1026.50	29.691756	21.84	22.455239
congestion_surcharge	0.00	2.75	2.176726	2.50	0.838856
airport_fee	0.00	1.75	0.137468	0.00	0.457917

```

df['pickup_hour'] = df['tpep_pickup_datetime'].dt.hour
df['pickup_day'] = df['tpep_pickup_datetime'].dt.day
df['pickup_weekend'] = df['tpep_pickup_datetime'].dt.weekday >= 4 # Friday=4
df['pickup_month'] = df['tpep_pickup_datetime'].dt.month
df['pickup_year'] = df['tpep_pickup_datetime'].dt.year
df['dropoff_hour'] = df['tpep_dropoff_datetime'].dt.hour
df['dropoff_day'] = df['tpep_dropoff_datetime'].dt.day
df['dropoff_weekend'] = df['tpep_dropoff_datetime'].dt.weekday >= 4 # Friday=4
df['dropoff_month'] = df['tpep_dropoff_datetime'].dt.month
df['dropoff_year'] = df['tpep_dropoff_datetime'].dt.year

df.head()

```

VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
131072	2023-01-12 09:58:04	2023-01-12 10:13:08	1.0
131073	2023-01-12 09:07:47	2023-01-12 09:36:28	1.0
131074	2023-01-12 09:39:29	2023-01-12 09:58:25	1.0
131075	2023-01-12 09:14:18	2023-01-12 09:29:54	1.0
131076	2023-01-12 09:41:56	2023-01-12 09:59:12	1.0

5 rows × 28 columns

```
df['trip_duration_min'] = (
    (df['tpep_dropoff_datetime'] - df['tpep_pickup_datetime'])
    .dt.total_seconds() / 60
)

df = df[df['trip_duration_min'] > 0]
```

```
df['fare_per_mile'] = df['fare_amount'] / df['trip_distance']
```

```
df['tip_percent'] = df['tip_amount'] / df['fare_amount']
```

```
df.groupby('pickup_weekend')['total_amount'].mean()
```

total_amount	
pickup_weekend	
False	29.972848
True	29.212195

dtype: float64

```
df.groupby('pickup_weekend')['total_amount'].median()
```

total_amount	
pickup_weekend	
False	22.2
True	21.3

dtype: float64

```
df['calculated_total'] = (
    df['fare_amount'] +
    df['extra'] +
    df['mta_tax'] +
```

```

        df['tip_amount'] +
        df['tolls_amount'] +
        df['improvement_surcharge'] +
        df['congestion_surcharge'] +
        df['airport_fee']
    )

df['total_diff'] = df['total_amount'] - df['calculated_total']

```

```
df['total_diff'].abs().describe().round(2)
```

total_diff	
count	1609793.00
mean	0.70
std	1.15
min	0.00
25%	0.00
50%	0.00
75%	2.50
max	4.50

dtype: float64

```
df.head()
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
131072	2	2023-01-12 09:58:04	2023-01-12 10:13:08	1.0
131073	2	2023-01-12 09:07:47	2023-01-12 09:36:28	1.0
131074	2	2023-01-12 09:39:29	2023-01-12 09:58:25	1.0
131075	2	2023-01-12 09:14:18	2023-01-12 09:29:54	1.0
131076	2	2023-01-12 09:41:56	2023-01-12 09:59:12	1.0

5 rows × 33 columns

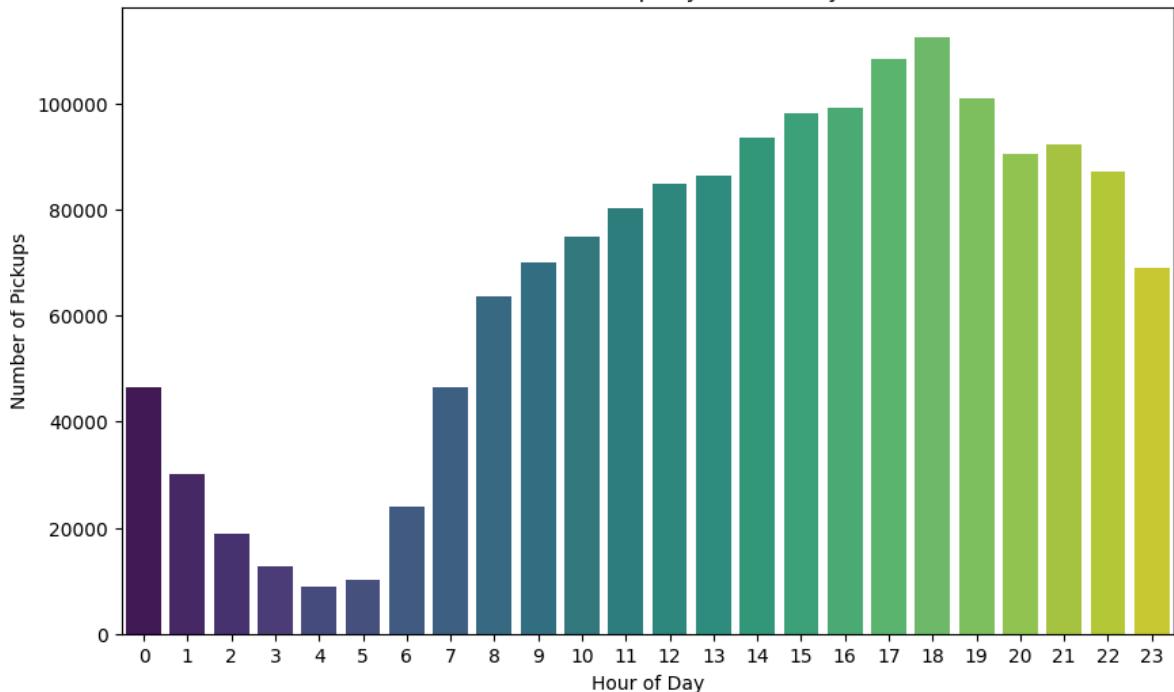
```

import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10, 6))
sns.countplot(x='pickup_hour', data=df, palette='viridis')
plt.title('Number of Pickups by Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('Number of Pickups')
plt.xticks(range(0, 24))
plt.show()

```

```
tmp/ipython-input-1569673394.py:5: FutureWarning:  
    assing `palette` without assigning `hue` is deprecated and will be removed in v  
    sns.countplot(x='pickup_hour', data=df, palette='viridis')  
    Number of Pickups by Hour of Day
```



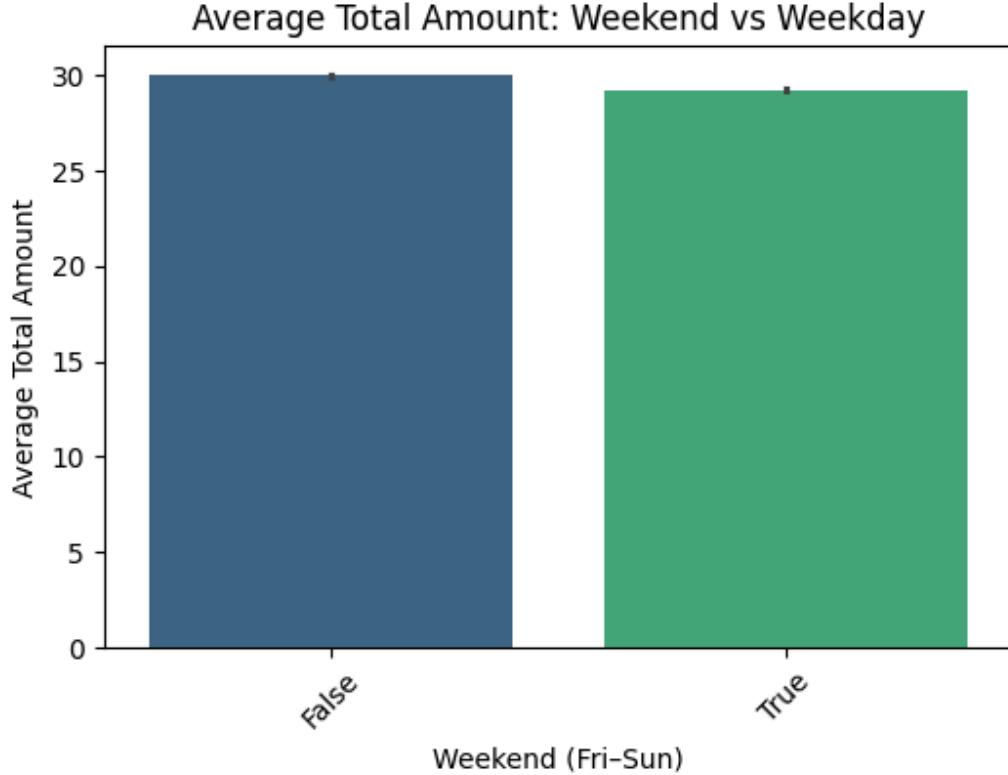
```
df.groupby('pickup_weekend')['total_amount'].mean().round(2)
```

total_amount	
pickup_weekend	
False	29.97
True	29.21

dtype: float64

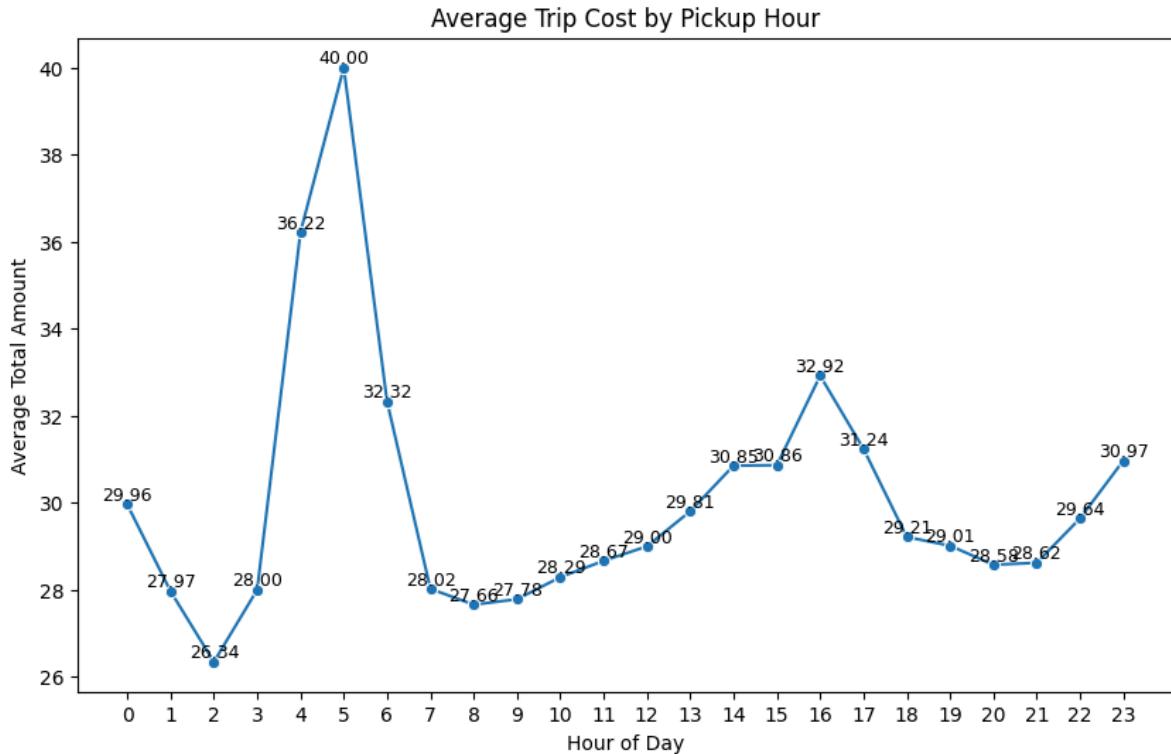
```
plt.figure(figsize=(6, 4))  
sns.barplot(x='pickup_weekend', y='total_amount', data=df, palette='viridis')  
plt.xticks(rotation=45)  
plt.title('Average Total Amount: Weekend vs Weekday')  
plt.xlabel('Weekend (Fri-Sun)')  
plt.ylabel('Average Total Amount')  
plt.show()
```

```
/tmp/ipython-input-1988818216.py:2: FutureWarning:  
  Passing `palette` without assigning `hue` is deprecated and will be removed in  
  sns.barplot(x='pickup_weekend', y='total_amount', data=df, palette='viridis')
```



```
hourly_avg = (  
    df.groupby('pickup_hour')['total_amount']  
    .mean()  
    .reset_index()  
)  
  
plt.figure(figsize=(10, 6))  
sns.lineplot(  
    x='pickup_hour',  
    y='total_amount',  
    data=hourly_avg,  
    marker='o'  
)  
  
for _, row in hourly_avg.iterrows():  
    plt.text(  
        row['pickup_hour'],  
        row['total_amount'],  
        f"{row['total_amount']:.2f}",  
        ha='center',  
        va='bottom',  
        fontsize=9  
)  
  
plt.title('Average Trip Cost by Pickup Hour')  
plt.xlabel('Hour of Day')
```

```
plt.ylabel('Average Total Amount')
plt.xticks(range(0, 24))
plt.show()
```



```
outliers = df[
    (df['trip_distance'] > 50) |
    (df['fare_amount'] > 300)
]

plt.figure(figsize=(8, 6))

# נקודות רגילות
sns.scatterplot(
    x='trip_distance',
    y='fare_amount',
    data=df,
    alpha=0.05,
    color='blue'
)

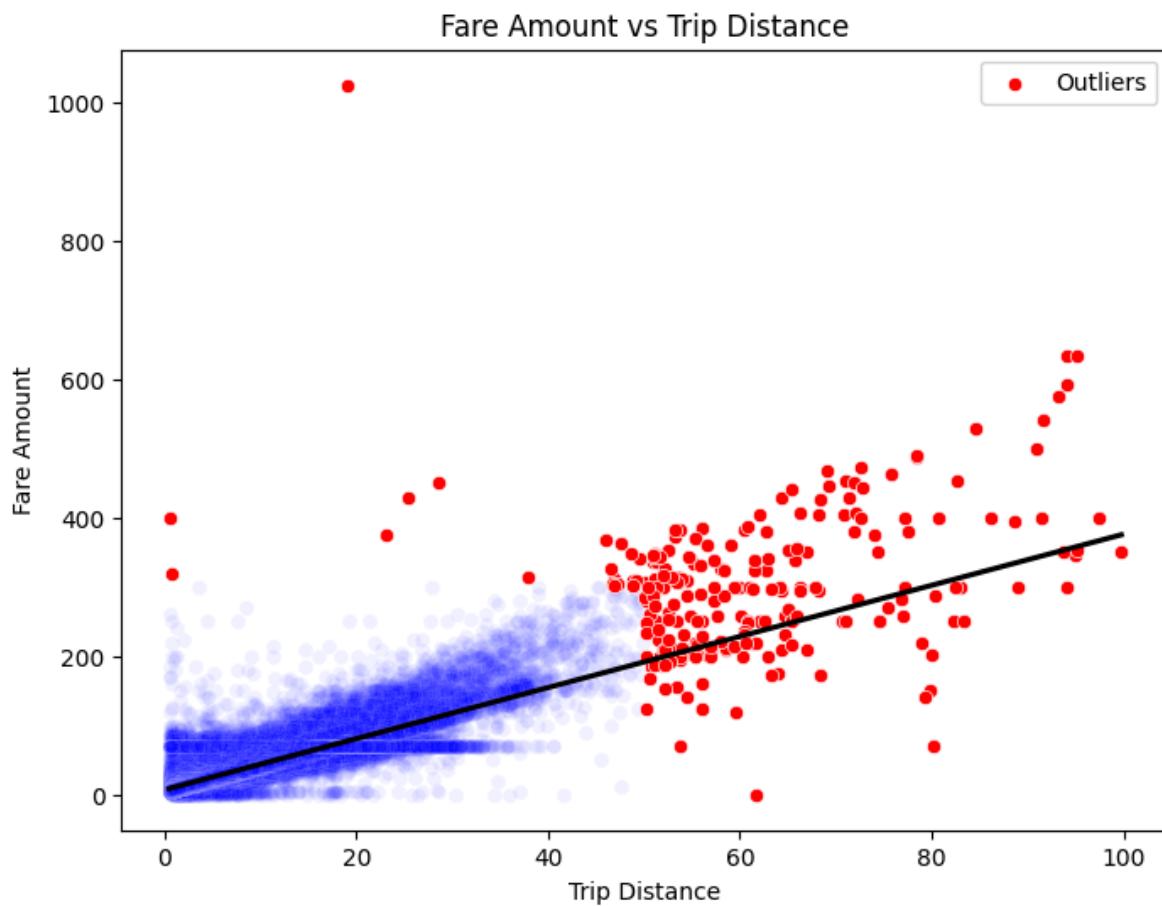
# חריגים
sns.scatterplot(
    x='trip_distance',
    y='fare_amount',
    data=outliers,
    color='red',
    label='Outliers'
)

# קו רגסיה
sns.regplot(
    x='trip_distance',
```

```
y='fare_amount',
data=df,
scatter=False,
color='black'
)

plt.title('Fare Amount vs Trip Distance')
plt.xlabel('Trip Distance')
plt.ylabel('Fare Amount')
plt.legend()
plt.show()
```

/usr/local/lib/python3.12/dist-packages/IPython/core/pylabtools.py:151: UserWarning:



```
payment_map = {
    0: 'Flex Fare trip',
    1: 'Credit card',
    2: 'Cash',
    3: 'No charge',
    4: 'Dispute',
    5: 'Unknown',
    6: 'Voided trip'
}
```

```
df['payment_type_label'] = df['payment_type'].map(payment_map)
```

```
plt.figure(figsize=(8, 5))

ax = sns.barplot(
    x='payment_type_label',
    y='tip_percent',
    data=df,
    palette='viridis'
)

# הוספת הערכיות מעל העמודות
for p in ax.patches:
    value = p.get_height()
    ax.annotate(
        f'{value:.2f}',
        (p.get_x() + p.get_width() / 2., value),
        ha='center',
        va='bottom',
        fontsize=10
    )

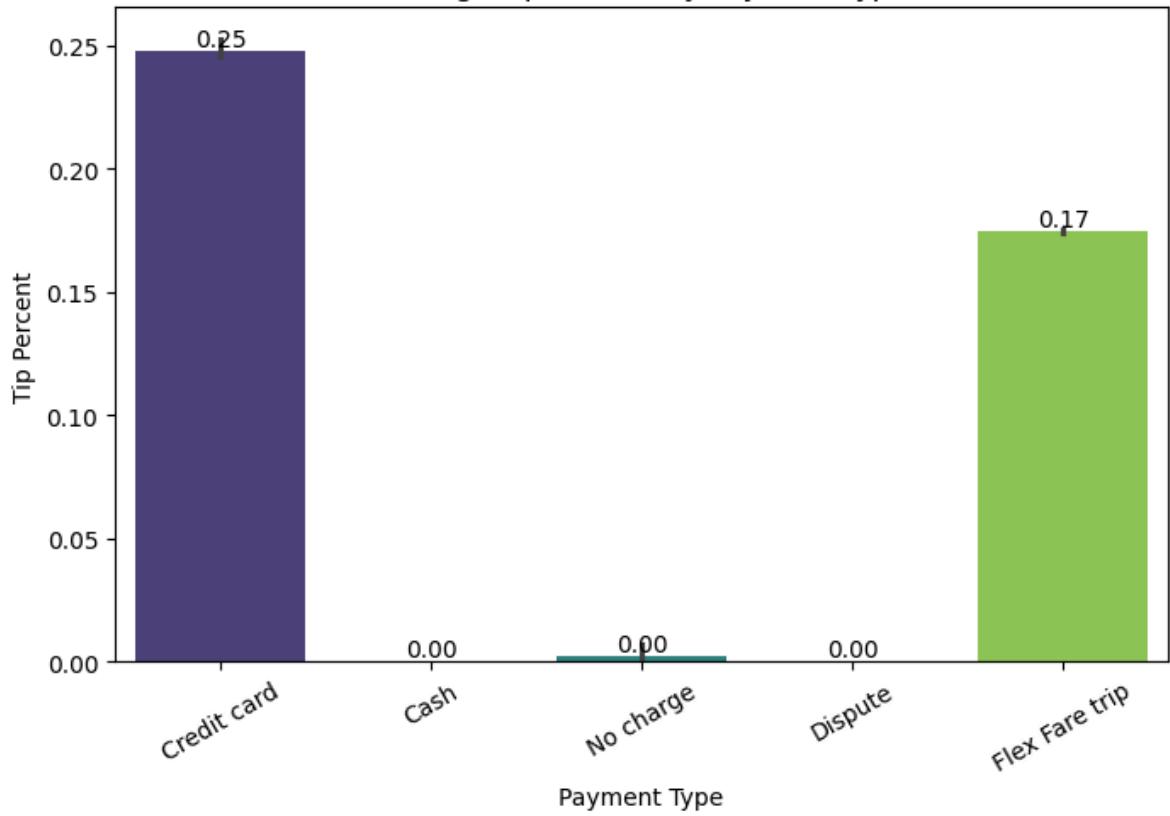
plt.title('Average Tip Percent by Payment Type')
plt.xlabel('Payment Type')
plt.ylabel('Tip Percent')
plt.xticks(rotation=30)
plt.show()
```

```
/tmp/ipython-input-1626809601.py:16: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in
```

```
ax = sns.barplot(
```

Average Tip Percent by Payment Type



```
ratecode_map = {
    1: 'Standard',
    2: 'JFK',
    3: 'Newark',
    4: 'Nassau/Westchester',
    5: 'Negotiated',
    6: 'Group ride',
    99: 'Unknown'
}

df['Ratecode_label'] = df['RatecodeID'].map(ratecode_map)

plt.figure(figsize=(8, 5))

ax = sns.barplot(
    x='Ratecode_label',
    y='total_amount',
    data=df,
    palette='viridis'
)

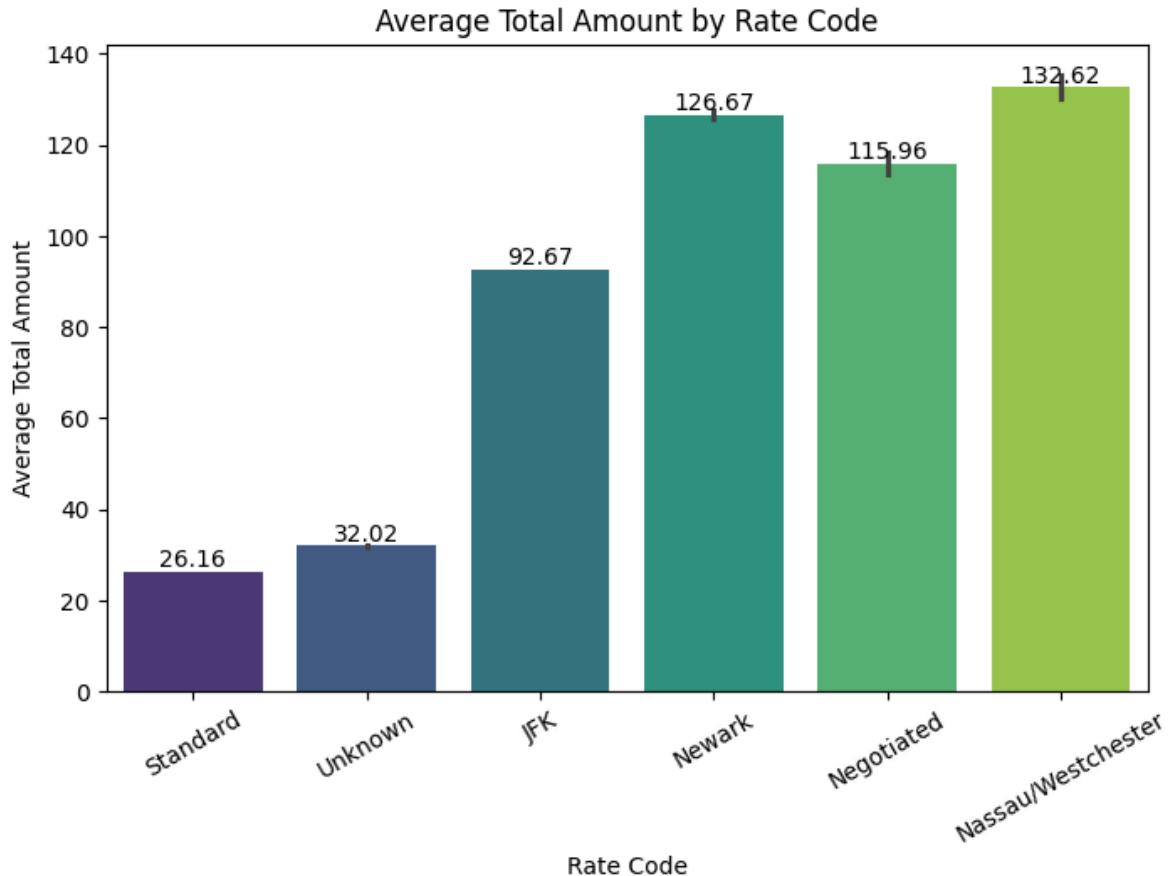
# הוספה הערכיהם מעל העמודות #
```

```
for p in ax.patches:  
    value = p.get_height()  
    ax.annotate(  
        f"{value:.2f}",  
        (p.get_x() + p.get_width() / 2., value),  
        ha='center',  
        va='bottom',  
        fontsize=10  
    )  
  
plt.title('Average Total Amount by Rate Code')  
plt.xlabel('Rate Code')  
plt.ylabel('Average Total Amount')  
plt.xticks(rotation=30)  
plt.show()
```

/tmp/ipython-input-2133174923.py:16: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.12.0

```
ax = sns.barplot(
```



```
df.groupby('Ratecode_label')['trip_distance'].mean().round(2)
```

trip_distance	
Ratecode_label	
JFK	18.25
Nassau/Westchester	20.46
Negotiated	15.50
Newark	18.06
Standard	2.89
Unknown	4.84

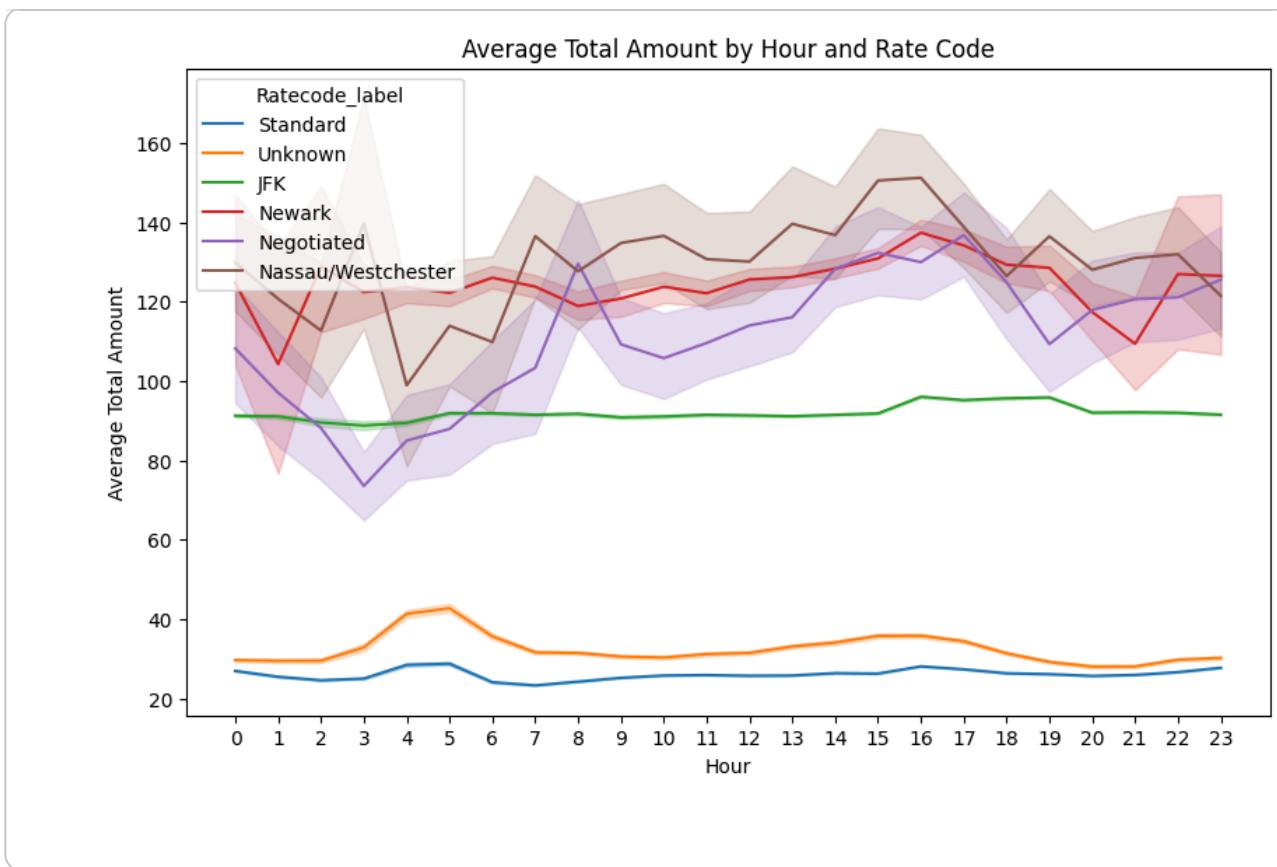
dtype: float64

```
df.groupby('Ratecode_label')['trip_duration_min'].mean().round(2)
```

trip_duration_min	
Ratecode_label	
JFK	52.79
Nassau/Westchester	42.49
Negotiated	40.80
Newark	44.31
Standard	16.41
Unknown	21.94

dtype: float64

```
plt.figure(figsize=(10, 6))
sns.lineplot(
    x='pickup_hour',
    y='total_amount',
    hue='Ratecode_label',
    data=df,
    estimator='mean'
)
plt.title('Average Total Amount by Hour and Rate Code')
plt.xlabel('Hour')
plt.ylabel('Average Total Amount')
plt.xticks(range(0, 24))
plt.show()
```



```
df.groupby(['Ratecode_label', 'pickup_weekend'])['total_amount'].mean().round(2)
```

total_amount		
Ratecode_label	pickup_weekend	
JFK	False	93.06
	True	92.05
Nassau/Westchester	False	131.29
	True	134.93
Negotiated	False	118.54
	True	112.37
Newark	False	127.22
	True	125.88
Standard	False	26.61
	True	25.37
Unknown	False	32.12
	True	31.89

dtype: float64

```
df[df['Ratecode_label'] == 'Unknown'].shape[0] / df.shape[0]
```

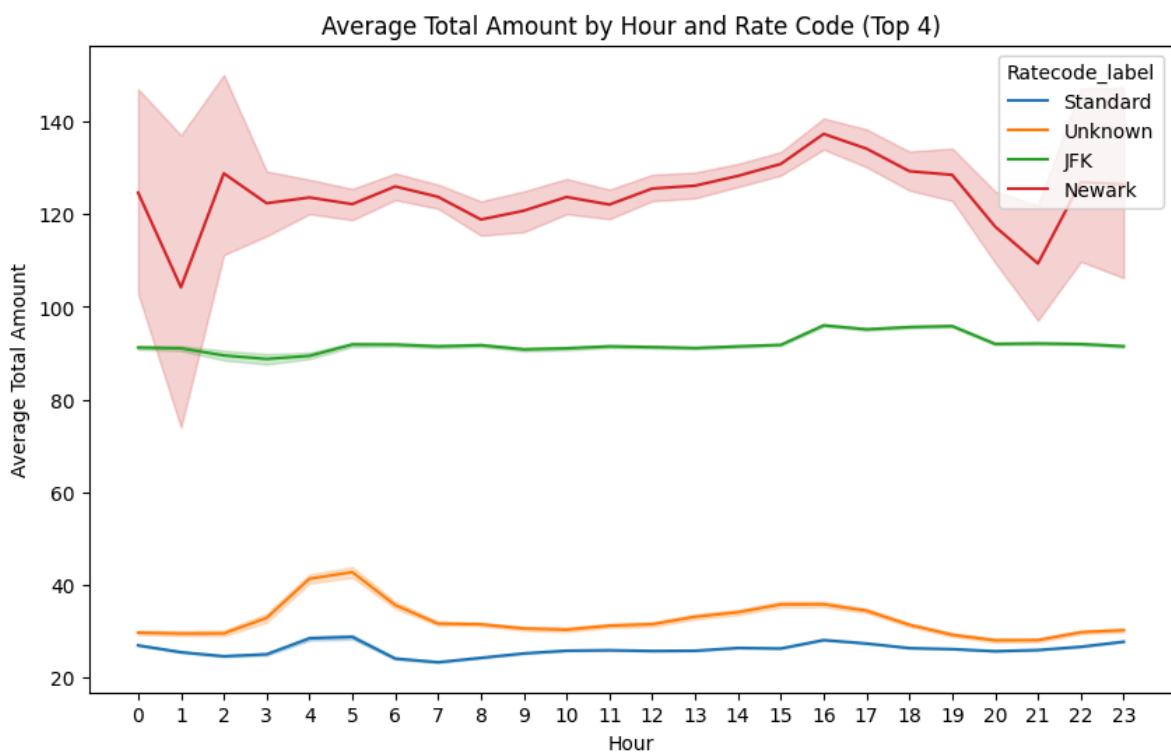
0 .06818764897101677

```
df.groupby('Ratecode_label')[['trip_distance','trip_duration_min','total_amount']]
```

Ratecode_label	trip_distance	trip_duration_min	total_amount
JFK	18.25	52.79	92.67
Nassau/Westchester	20.46	42.49	132.62
Negotiated	15.50	40.80	115.96
Newark	18.06	44.31	126.67
Standard	2.89	16.41	26.16
Unknown	4.84	21.94	32.02

```
top_ratecodes = df['Ratecode_label'].value_counts().head(4).index # 4 最高のラートコード
df_top = df[df['Ratecode_label'].isin(top_ratecodes)]

plt.figure(figsize=(10,6))
sns.lineplot(x='pickup_hour', y='total_amount', hue='Ratecode_label', data=df_top)
plt.title('Average Total Amount by Hour and Rate Code (Top 4)')
plt.xlabel('Hour')
plt.ylabel('Average Total Amount')
plt.xticks(range(0,24))
plt.show()
```



```
df['total_diff_abs'] = df['total_diff'].abs()  
df['total_diff_abs'].describe().round(2)
```

total_diff_abs	
count	1609793.00
mean	0.70
std	1.15
min	0.00
25%	0.00
50%	0.00
75%	2.50
max	4.50

dtype: float64

```
{  
    "diff<=0.01": (df['total_diff_abs'] <= 0.01).mean(),  
    "diff<=1.00": (df['total_diff_abs'] <= 1.00).mean(),  
    "diff<=2.50": (df['total_diff_abs'] <= 2.50).mean(),  
}  
  
{'diff<=0.01': np.float64(0.7225450725652305),  
 'diff<=1.00': np.float64(0.7225463149609919),  
 'diff<=2.50': np.float64(0.9622821070783635)}
```

```
df['has_tip'] = df['tip_amount'] > 0
```

```
df['has_tip'].mean().round(3) * 100 # %  
np.float64(79.2)
```

```
df.groupby('payment_type_label')['has_tip'].agg(  
    trips='count',  
    tip_rate='mean'  
) .round(3)
```

	trips	tip_rate
payment_type_label		
Cash	250384	0.000
Credit card	1244981	0.956
Dispute	8639	0.001
Flex Fare trip	101554	0.823
No charge	4235	0.000

```
plt.figure(figsize=(8,5))
ax = sns.barplot(
    x='payment_type_label',
    y='has_tip',
    data=df,
    palette='viridis'
)

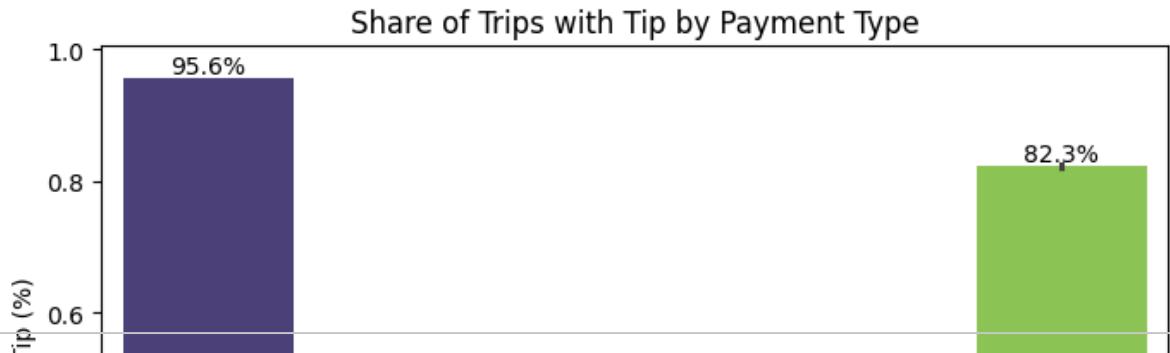
for p in ax.patches:
    v = p.get_height()
    ax.annotate(
        f'{v*100:.1f}%',
        (p.get_x() + p.get_width()/2, v),
        ha='center',
        va='bottom',
        fontsize=10
    )

plt.title('Share of Trips with Tip by Payment Type')
plt.xlabel('Payment Type')
plt.ylabel('Trips with Tip (%)')
plt.xticks(rotation=30)
plt.show()
```

```
/tmp/ipython-input-1545293828.py:2: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v

```
ax = sns.barplot(
```



```
df.groupby('payment_type_label').agg(
    tip_rate=('has_tip', 'mean'),
    avg_tip_amount=('tip_amount', 'mean'),
    avg_tip_percent=('tip_percent', 'mean'),
    trips=('has_tip', 'count')
).round(3)
```

payment_type_label	tip_rate	avg_tip_amount	avg_tip_percent	trips
Cash	0.000	0.000	0.0%	
Credit card	0.956	4.462	inf	1244981
Dispute	0.001	0.003	0.00%	8639
Flex Fare trip	0.823	3.770	inf	101554
No charge	0.000	0.024	0.003	4235

```
df.groupby('pickup_hour')['has_tip'].mean().round(3)
```

has_tip	
pickup_hour	
<hr/>	
0	0.789
1	0.782
2	0.771
3	0.749
4	0.704
5	0.698
6	0.748
7	0.806
8	0.818
9	0.790
10	0.768
11	0.764
12	0.766
13	0.762
14	0.765
15	0.772
16	0.778
17	0.802
18	0.816
19	0.818
20	0.822
21	0.827
22	0.822
23	0.807

dtype: float64

```
df.groupby('pickup_weekend')['has_tip'].mean().round(3)
```

```
has_tip
```

```
pickup_weekend
```

False	0.800
True	0.778

```
dtype: float64
```

```
df.groupby('Ratecode_label')['has_tip'].mean().round(3)
```

```
has_tip
```

```
Ratecode_label
```

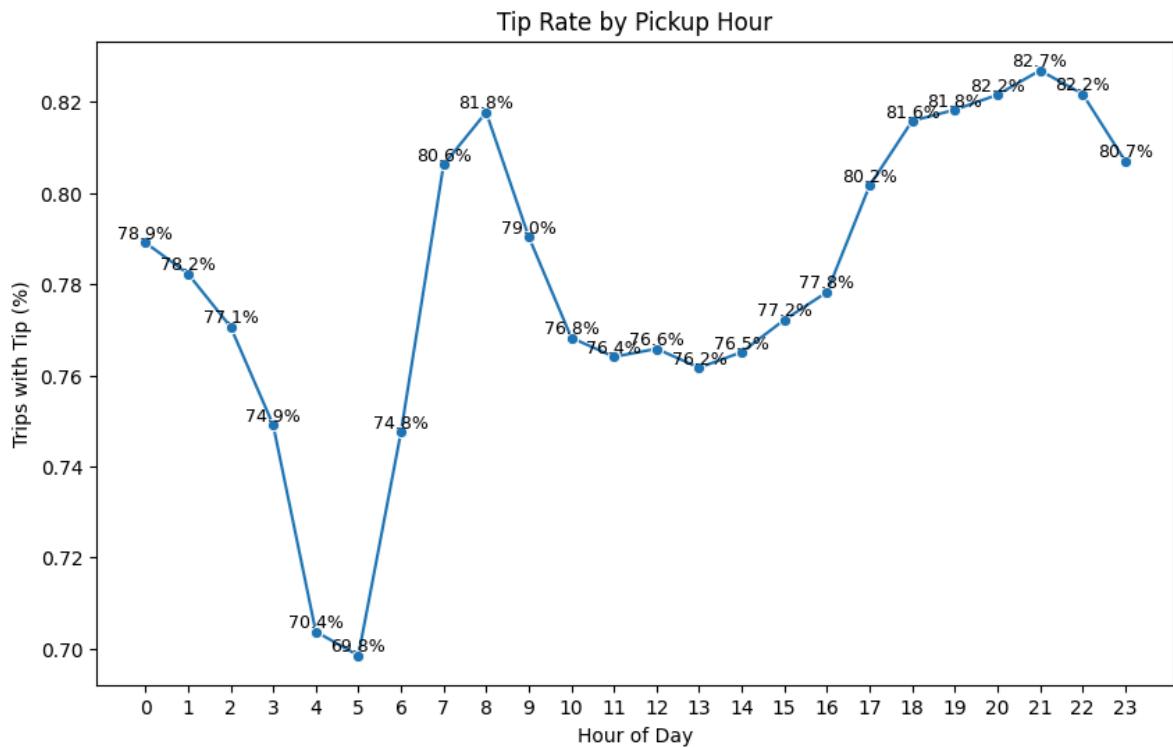
JFK	0.761
Nassau/Westchester	0.614
Negotiated	0.663
Newark	0.753
Standard	0.796
Unknown	0.762

```
dtype: float64
```

```
plt.figure(figsize=(10,6))
sns.lineplot(x='pickup_hour', y='has_tip', data=tip_by_hour, marker='o')

for _, row in tip_by_hour.iterrows():
    plt.text(
        row['pickup_hour'],
        row['has_tip'],
        f'{row["has_tip"]*100:.1f}%',
        ha='center',
        va='bottom',
        fontsize=9
    )

plt.title('Tip Rate by Pickup Hour')
plt.xlabel('Hour of Day')
plt.ylabel('Trips with Tip (%)')
plt.xticks(range(0,24))
plt.show()
```



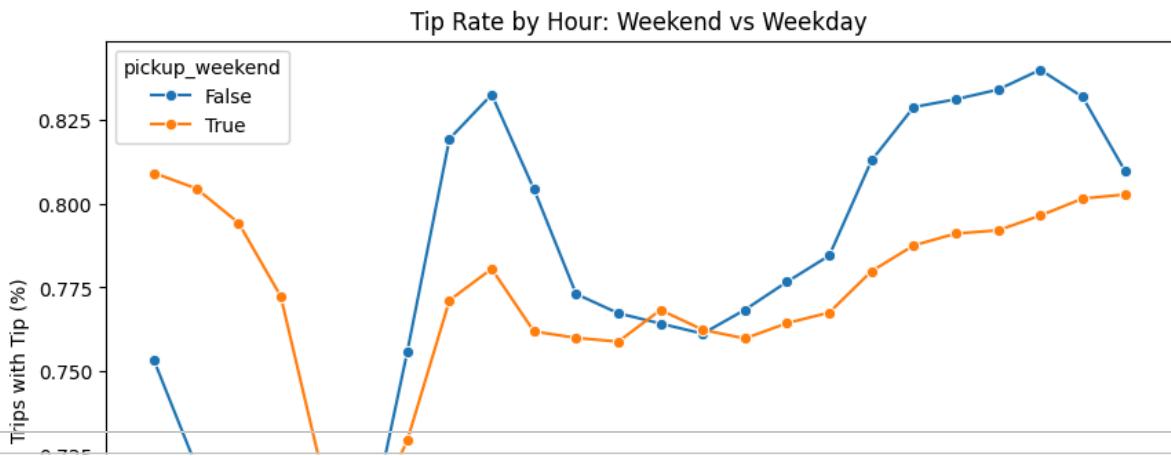
```

tip_hour_weekend = (
    df.groupby(['pickup_hour','pickup_weekend'])['has_tip']
    .mean()
    .reset_index()
)

plt.figure(figsize=(10,6))
sns.lineplot(
    x='pickup_hour',
    y='has_tip',
    hue='pickup_weekend',
    data=tip_hour_weekend,
    marker='o'
)

plt.title('Tip Rate by Hour: Weekend vs Weekday')
plt.xlabel('Hour of Day')
plt.ylabel('Trips with Tip (%)')
plt.xticks(range(0,24))
plt.show()

```



	pickup_hour	trips	avg_fare	tip_rate
0	0	46517	29.960	0.789
1	1	30063	27.966	0.782
2	2	18875	26.344	0.771
3	3	12702	28.000	0.749
4	4	8985	36.222	0.704
5	5	10092	40.000	0.698
6	6	24055	32.324	0.748
7	7	46423	28.023	0.806
8	8	63742	27.660	0.818
9	9	69999	27.783	0.790
10	10	74901	28.291	0.768
11	11	80214	28.666	0.764
12	12	84896	28.999	0.766
13	13	86444	29.805	0.762
14	14	93597	30.854	0.765
15	15	98230	30.864	0.772
16	16	99100	32.921	0.778
17	17	108494	31.240	0.802
18	18	112565	29.214	0.816
19	19	101032	29.013	0.818
20	20	90538	28.576	0.822
21	21	92235	28.620	0.827
22	22	87121	29.640	0.822
23	23	68973	30.967	0.807

```

driver_weekend = (
    df.groupby(['pickup_hour','pickup_weekend'])
    .agg(
        trips=('total_amount','count'),
        avg_fare=('total_amount','mean'),
        tip_rate=('has_tip','mean')
    )
    .round(3)
    .reset_index()
)

driver_weekend

```



```
df['fare_per_mile'] = df['fare_amount'] / df['trip_distance']

suspicious_short_expensive = df[
    (df['trip_distance'] < 1) &
    (df['fare_amount'] > df['fare_amount'].quantile(0.99))
]

suspicious_short_expensive[['trip_distance', 'fare_amount', 'fare_per_mile']].de
```

	trip_distance	fare_amount	fare_per_mile
count	51.00	51.00	51.00
mean	0.70	117.51	175.11
std	0.12	67.08	112.60
min	0.51	78.00	96.67
25%	0.60	84.15	111.92
50%	0.70	87.30	131.62
75%	0.80	110.00	174.07
max	0.96	400.00	666.67

```
threshold = df['fare_per_mile'].quantile(0.995)

suspicious_fpm = df[df['fare_per_mile'] > threshold]
```

```
suspicious_fpm[['trip_distance', 'fare_amount', 'fare_per_mile']]
```

1	trip_distance	fare_amount	fare_per_mile	
131785	0.60	15.5	25.833333	0.722
131873	1	4.00	True 82.0	26.363 20.500000 0.805
132002	0.88	17.0	19.318182	0.688
132537	2	0.79	True 15.6	25.241 19.746835 0.794
133493	0.81	17.0	20.987654	0.667
7 ...	3	...	True 9891	26.684 ... 0.772
4256995	1.60	70.0	43.750000	0.675
4257067	4	0.75	True 17.7	23.415 23.600000 0.719
4257077	0.54	10.7	19.814815	0.696
4257192	5	0.66	True 12.8	41.932 19.393939 0.702
4258617	0.73	65.0	89.041096	0.756
8048 rows × 3 columns		True 7557	35.914	0.729
14	7	False 33767	27.009	0.819

```
suspicious_fast = df[
    (df['trip_duration_min'] < 5) &
    (df['total_amount'] > df['total_amount'].quantile(0.99))
]
```

```
suspicious_fast[['trip_duration_min', 'total_amount']].describe().round(2)
```

20	10	False	47297	28.493	0.773
21	trip_duration_min	total_amount			
	10	True	27604	27.946	0.760
count			71.00	71.00	
--	..		71.00	71.00	29.200 0.767
mean			1.50	152.83	
22	11	True	30827	27.810	0.759
std			1.44	68.87	
--	12		1.44	1.44	29.421 0.764
min			0.03	104.33	
25	12	True	33098	28.337	0.768
25%			0.38	107.33	
--	13		0.38	107.33	30.241 0.761
50%			0.80	121.10	
27	13	True	34218	29.141	0.762
75%			2.41	181.62	
--	14		2.41	181.62	31.240 0.768
max			4.75	481.20	
29	14	True	35600	30.224	0.760

```
from sklearn.linear_model import LinearRegression
```

```
X = df[['trip_distance']]
y = df['fare_amount']
```

```

model = LinearRegression()
model.fit(X, y)

df['fare_pred'] = model.predict(X)
df['fare_residual'] = df['fare_amount'] - df['fare_pred']

# חירגים חיובים קיצוניים
suspicious_residuals = df[df['fare_residual'] > df['fare_residual'].quantile(0.7)]
suspicious_residuals

```

39	19	True	32402	28.151	0.791	
40	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count		
	20	False	63671	28.314	0.834	
131873	1	2023-01-12 10:45:48		2023-01-12 10:46:03		1.0
132088	21	2023-01-12 11:44:10		2023-01-12 12:22:20		1.0
132240	2	2023-01-12 11:37:09		2023-01-12 12:12:50		2.0
132509	22	2023-01-12 11:32:36		2023-01-12 12:17:04		1.0
132546	2	2023-01-12 11:36:32		2023-01-12 11:59:56		1.0
46	...	False	41675	31.244	0.810	...
4258817	2	2023-11-09 23:48:33		2023-11-10 01:15:19		1.0
4259354	2	2023-11-09 23:52:51		2023-11-10 00:27:32		1.0
4259357	2	2023-11-09 23:23:11		2023-11-09 23:55:39		1.0
4259442	2	2023-11-10 00:29:42		2023-11-10 01:11:39		1.0
4259783	2	2023-11-10 00:17:23		2023-11-10 00:54:16		1.0

8049 rows × 39 columns

```

suspicious_tips = df[
    (df['tip_percent'] > 0.6) &
    (df['total_amount'] < 20)
]

suspicious_tips

```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
131555	2	2023-01-12 10:35:53	2023-01-12 10:41:02	2.0
133487	1	2023-01-12 13:34:53	2023-01-12 13:41:24	2.0
133731	1	2023-01-12 13:57:45	2023-01-12 14:06:08	1.0
134574	1	2023-01-12 14:14:14	2023-01-12 14:21:26	1.0

```
df[
    (df['pickup_hour'].between(3,5)) &
    (df['fare_per_mile'] > df['fare_per_mile'].quantile(0.99))
]
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
158774	2	2023-01-14 04:33:29	2023-01-14 04:40:09	2.0
299529	2	2023-01-28 03:48:36	2023-01-28 03:53:56	1.0
299535	1	2023-01-25 05:19:37	2023-01-25 05:27:06	1.0
313368	2	2023-01-26 05:43:57	2023-01-26 05:44:13	1.0
378428	2	2023-01-31 04:59:43	2023-01-31 05:02:19	1.0
...
4122368	1	2023-10-29 04:02:18	2023-10-29 04:09:14	1.0
4122384	2	2023-10-29 04:19:31	2023-10-29 04:30:45	1.0
4124611	2	2023-11-01 03:08:57	2023-11-01 03:14:00	1.0
4124760	1	2023-11-01 04:01:03	2023-11-01 04:29:48	1.0
4243767	1	2023-11-09 05:22:57	2023-11-09 05:48:53	1.0

143 rows × 39 columns